

Author Attribution

Aleena Syed

CS 4395.001

```
In [71]: import pandas as pd
        from sklearn.model_selection import train_test_split
        from nltk.corpus import stopwords
        from sklearn.feature_extraction.text import TfidfVectorizer
        from sklearn.naive_bayes import BernoulliNB
        from sklearn.linear_model import LogisticRegression
        from sklearn.metrics import accuracy_score
        from sklearn.neural_network import MLPClassifier
        from sklearn.neural_network import MLPRegressor

In [72]: # Convert author column to categorical data
df = pd.read_csv('federalist.csv')
df['author'] = df['author'].astype('category')
print(df.head())

      author      text
0  HAMILTON  FEDERALIST. No. 1 General Introduction For the...
1      JAY    FEDERALIST No. 2 Concerning Dangers from Forei...
2      JAY    FEDERALIST No. 3 The Same Subject Continued (C...
3      JAY    FEDERALIST No. 4 The Same Subject Continued (C...
4      JAY    FEDERALIST No. 5 The Same Subject Continued (C...

In [73]: # display counts by author
df['author'].value_counts()

Out[73]: HAMILTON      49
MADISON      15
HAMILTON OR MADISON  11
JAY           5
HAMILTON AND MADISON  3
Name: author, dtype: int64

In [74]: # divide into train and test, with 80% in train, using random state 1234
X = df.text
y = df.author

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, train_size=0.8, random_state=1234)

# print shape of train and test
print(X_train.shape)
print(X_test.shape)

(66,)
(17,)

In [75]: # remove stopwords and perform tf-idf vectorization
stopwords = set(stopwords.words('english'))
vectorizer = TfidfVectorizer(stop_words=stopwords)

X_train = vectorizer.fit_transform(X_train)
X_test = vectorizer.transform(X_test)

# print shape of train and test
print(X_train.shape)
print(X_test.shape)

(66, 7876)
(17, 7876)

In [76]: # Bernoulli Naive Bayes model
naive_bayes = BernoulliNB()
naive_bayes.fit(X_train, y_train)

Out[76]: BernoulliNB()

In [77]: # accuracy on test
pred = naive_bayes.predict(X_test)
print('accuracy score: ', accuracy_score(y_test, pred))

accuracy score:  0.5882352941176471

In [78]: # second Bernoulli Naive Bayes model
X = df.text
y = df.author
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, train_size=0.8, random_state=1234)

In [79]: # max_features = 1000, add bigrams as feature
vectorizer2 = TfidfVectorizer(stop_words=stopwords, max_features=1000, ngram_range=(1,2))
X_train = vectorizer2.fit_transform(X_train)
X_test = vectorizer2.transform(X_test)

In [80]: # try new model with updated parameters
naive_bayes = BernoulliNB()
naive_bayes.fit(X_train, y_train)

Out[80]: BernoulliNB()

In [81]: # new accuracy on test
pred = naive_bayes.predict(X_test)
print('accuracy score: ', accuracy_score(y_test, pred))

accuracy score:  0.9411764705882353

In [82]: # logistic regression
X = df.text
y = df.author

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, train_size=0.8, random_state=1234)

vectorizer = TfidfVectorizer(binary=True)
X_train = vectorizer.fit_transform(X_train)
X_test = vectorizer.transform(X_test)

classifier = LogisticRegression(C = 0.1, solver='lbfgs', class_weight='balanced')
classifier.fit(X_train, y_train)

# accuracy on test
pred = classifier.predict(X_test)
print('accuracy score: ', accuracy_score(y_test, pred))

accuracy score:  0.7647058823529411

In [83]: # neural networks topology 1
vectorizer = TfidfVectorizer(stop_words=stopwords, binary=True)
X = vectorizer.fit_transform(df.text)
y = df.author

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, train_size=0.8, random_state=1234)

classifier = MLPClassifier(solver='lbfgs', alpha=1e-5,
                           hidden_layer_sizes=(15, 2), random_state=1)
classifier.fit(X_train, y_train)
pred = classifier.predict(X_test)
print('accuracy score: ', accuracy_score(y_test, pred))

accuracy score:  0.7647058823529411

In [84]: # neural networks topology 2
vectorizer = TfidfVectorizer(stop_words=stopwords, binary=True)
X = vectorizer.fit_transform(df.text)
y = df.author

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, train_size=0.8, random_state=1234)

classifier = MLPClassifier(alpha=1e-05, hidden_layer_sizes=(15,), random_state=1,
                           solver='lbfgs')
classifier.fit(X_train, y_train)
pred = classifier.predict(X_test)
print('accuracy score: ', accuracy_score(y_test, pred))

accuracy score:  0.7647058823529411

In [88]: # neural networks topology 3
vectorizer = TfidfVectorizer(stop_words=stopwords, binary=True)
X = vectorizer.fit_transform(df.text)
y = df.author

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, train_size=0.8, random_state=1234)

classifier = MLPClassifier(hidden_layer_sizes=(8,8,8), activation='relu', solver='adam', max_iter=1000)

classifier.fit(X_train, y_train)
pred = classifier.predict(X_test)
print('accuracy score: ', accuracy_score(y_test, pred))

accuracy score:  0.5882352941176471
```

multi-layer perceptron using backpropagation gives best accuracy